

Model Catalog

Explore 67 models available in the MindRouter cluster

All

LLM

Embedding

Reranking

All

Ollama

vLLM

Vision

Thinking

Tools

**Nemotron-3-Super-120b**
nvidia · vllm

NVIDIA Nemotron-3-Super-120b is an open hybrid Mixture-of-...

Thinking LLM

Context Length: 256k

**Qwen/Qwen3-32B**
qwen · vllm

Qwen3-32B is a 32-billion-parameter dense...

Thinking Tools LLM

Context Length: 32k

**Qwen/Qwen3-Embedding-8B**
qwen · vllm

Qwen3-Embedding-8B is a transformer-based embedding...

Tools Embedding Embedding

Context Length: 8k

**Qwen/Qwen3-Reranker-8B**
qwen · vllm

Qwen3-Reranker-8B is an 8-billion-parameter transforme...

Tools Reranking

Context Length: 8k

**aya-expanses:32b**
command · ollama

Aya Expanses 32B is a transformer-based multilingual large language model released by Cohere Labs. With 32 billion parameters it targets a wide range of languages (over 20) and

**aya-expanses:8b**
command · ollama

Cohere Aya Expanses 8B is a multilingual large language model optimized for broad language coverage across 23+ languages. It delivers strong multilingual performance in a

**aya:35b**
command · ollama

Aya 35B is an instruction-tuned, open-access multilingual language model from Cohere Labs. Built on a transformer architecture, it extends the Aya family's focus on high-quality

**aya:8b**
command · ollama

Aya 8B is an open-access, multilingual language model from Cohere Labs that builds on the Aya instruction-tuned family. With 8 billion parameters and F16 precision, it is optimized for

**black-forest-labs/FLUX.2-klein-9B**
· diffusion

FLUX.2-dev is Black Forest Labs' next-generation text-to-image diffusion model, designed for high-fidelity photorealistic generation up to 4MP. It

**codeqwen:7b**
qwen · ollama

codeqwen:7b is a 7-billion-parameter, decoder-only transformer model from the Qwen 2.5-Coder family, fine-tuned for programming tasks. Optimized for a 32K token



codestral:22b
mistral · ollama

Codestral 22B is Mistral AI's first dedicated code-generation...

LLM

Context Length: 32k



command-r-plus:104b
command · ollama

Command R+ is Cohere's 104-billion-parameter...

Tools LLM

Context Length: 32k



command-r:35b
command · ollama

Cohere Command R 35B is a generative model optimized for...

Tools LLM

Context Length: 32k



dots.OCR
qwen · vllm

dots.OCR is a compact vision-language model built on ...

Vision LLM

Context Length: 32k



gemma3:12b
gemma · ollama

Gemma 3 12B is a 12-billion-parameter...

Vision LLM

Context Length: 32k



gemma3:1b
gemma · ollama

Gemma 3 1B is a compact member of Google's Gemma 3...

LLM

Context Length: 32k



gemma3:4b
gemma · ollama

Gemma 3 4B is a lightweight, multimodal language model fro...

Vision LLM

Context Length: 32k



google/gemma-4-26b
gemma · vllm

Gemma-4 26B is a Mixture-of-Experts (MoE) language model from Google DeepMind, built with a total of 26 billion parameters while activating only about 3.8 billion



google/gemma-4-31b
gemma · vllm

Gemma-4-31B is a 31-billion-parameter dense transformer from Google DeepMind's Gemma-4 series. It offers an exceptionally long 262k-token context window,



granite-code:8b
granite · ollama

Granite-Code 8B is a decoder-only, open-source...

LLM

Context Length: 32k

Mixture-of-Experts (MoE) model



granite3-dense:8b
granite · ollama

Granite 3 dense 8B is a decoder-only dense transformer model optimized for high-throughput inference with 4-bit quantization (Q4_K_M). It offers a 4096-token context



granite3-moe:3b
granite · ollama

Granite 3-MoE 3B is a 3-billion-parameter transformer from IBM's Granite 3 series that employs a Mixture-of-Experts (MoE) design to achieve low-latency inference. Optimized and instruction following.



leanstral:119b
mistral · ollama

- **Architecture/Family:** Built on the Mistral transformer family, adapted for the Lean 4 formal language [2] - **Parameter Count:** 119 billion parameters



llama3.1:8b
llama · ollama

Llama 3.1 8B is a lightweight member of Meta's Llama 3.1 family, offering 8 billion parameters and optimized for instruction-following tasks. Built

signature by the model name)

with the latest knowledge -

 **llama3.2-vision:11b** 
llama · ollama

Meta Llama 3.2 Vision 11B is a multimodal model that extends...

Vision **LLM**

Context Length: 32k

 **llama3.2-vision:90b** 
llama · ollama

Llama 3.2 Vision 90B is a multimodal large language mod...

Vision **LLM**

Context Length: 32k

 **llama3.2:1b** 
llama · ollama

Llama 3.2 1B is a lightweight, 1-billion-parameter member of...

Tools **LLM**

Context Length: 32k

 **llama3.2:3b** 
llama · ollama

Llama 3.2 3B is a 3-billion-parameter,...

Tools **LLM**

Context Length: 32k

 **llama3.3:70b** 
llama · ollama

Llama 3.3 70B is a multilingual, instruction-tuned generative...

Tools **LLM**

Context Length: 32k

 **llava-llama3:8b** 
llama · ollama

LLaVA-Llama3 8B is a multimodal chatbot model that merges...

Vision **LLM**

Context Length: 4k

 **llava:13b** 
llama · ollama

LLaVA-13B is a multimodal large language model that combines ...

Vision **LLM**

Context Length: 4k

 **llava:34b** 
llama · ollama

LLaVA-34B is a multimodal instruction-tuned model that...

Vision **LLM**

Context Length: 4k

 **llava:7b** 
llama · ollama

LLaVA-7B is a multimodal model that fuses a CLIP ViT-L/14 visual encoder with the Vicuna LLM (itself built on Meta's LLaMA-2) via a lightweight projection layer. Designed for vision-language

 **magistral:24b** 
mistral · ollama

Magistral 24B is a reasoning-optimized large language model built on the Mistral Medium 3 architecture. It targets complex, multi-step problem solving and

 **marco-o1:7b** 
· ollama

Marco-o17B is an open-source large-language model built on the Qwen2-7B-Instruct backbone and fine-tuned with chain-of-thought and instruction datasets to boost reasoning on Vision / LLaMA 2 based

 **microsoft/phi-4**
phi · vllm

The Phi-4 series from Microsoft is a family of compact, high-efficiency language models that extend the capabilities of the earlier Phi-3 line with enhanced reasoning and


 **mistral-large:123b** 
mistral · ollama

Mistral-large:123b is a flagship model from the Mistral 3 family, built on a transformer

 **mistral-nemo:12b** 
mistral · ollama

Mistral-NeMo 12B is a state-of-the-art large language model jointly developed by

architecture with a sparse
mixture-of-experts (MoE) design...




Mistral-Small:22B is a
22-billion-parameter transform...

Tools LLM

Context Length: 32k


Mistral AI and NVIDIA. Built on
the Mistral-Nemo architecture...



Mistral 7B is a
7-billion-parameter decoder-on...

Tools LLM


Context Length: 32k



Mixtral 8x22B is a
Mixture-of-Experts transformer...

Tools LLM


Context Length: 32k



Mixtral 8x7B is a
Mixture-of-Experts language...

LLM


Context Length: 32k



Nemotron-Mini-4B is a compact,
distilled variant of NVIDIA's...

LLM


Context Length: 4k



Nemotron-70B is a
70-billion-parameter...

LLM


Context Length: 32k



Olmo 2 13B is a fully open,
decoder-only transformer from...

LLM


Context Length: 4k




A 120-billion-parameter
open-weight model from
OpenAI's GPT-OSS family,
gpt-oss-120b is built for
production-grade,
general-purpose workloads that




The openai/gpt-oss-20b model is
a 20-billion-parameter
mixture-of-experts language
model released under an
Apache 2.0 license. It is
optimized for low-latency,




Phi-3 14B is a
14-billion-parameter
decoder-only transformer from
Microsoft's Phi-3 family,
optimized for fast inference on
consumer-grade hardware. It is




Phi-4-Reasoning is Microsoft's
14-billion-parameter small



Phi-3.5-mini is a
3.8 billion-parameter, lightweight
language model from Microsoft's
Phi-3 family. Built for multilingual
text generation, it is optimized
for low-latency and high



Phi-3 3.8B is a compact,
state-of-the-art language model
from Microsoft's Phi-3 family.
Built on a transformer
architecture with 3.8 billion
parameters, it is optimized for
language generation.



Qwen 3.5 122B is a mixture-of-
experts model from Alibaba with

language model, fine-tuned from
Phi-4 specifically for complex...

LLM ing Tools LLM

 **qwen/qwen3.5-35b**
qwen · vllm

Qwen 3.5-35B is a decoder-only
transformer from Alibaba Cloud...

Vision Thinking Tools LLM

Context Length: 128k

122B total parameters and
approximately 10B active per...

Vision Thinking Tools LLM

 **qwen/qwen3.6-27b** 
qwen · vllm

Qwen 3.6-27B is a
27-billion-parameter transform...

Vision Thinking Tools LLM

Context Length: 256k

 **qwen/qwen3.6-35b** 
qwen · vllm

Qwen 3.6-35B is a
35-billion-parameter...

Vision Thinking Tools LLM

Context Length: 256k

 **qwen2.5-coder:14b** 
qwen · ollama

Qwen2.5-Coder 14B is a
decoder-only transformer mod...

Tools LLM

Context Length: 32k

 **qwen2.5-coder:32b** 
qwen · ollama

Qwen2.5-Coder 32B is an
open-source, instruction-tuned...

Tools LLM

Context Length: 32k

 **qwen2.5-coder:7b** 
qwen · ollama

Qwen2.5-Coder 7B is a
7-billion-parameter,...

Tools LLM

Context Length: 32k

 **qwen2.5:0.5b** 
qwen · ollama

Qwen2.5 0.5B is a compact,
decoder-only transformer from...

Tools LLM

Context Length: 32k

 **qwen2.5:1.5b** 
qwen · ollama

Qwen2.5 1.5B is a compact,
1.5-billion-parameter transformer
model from the Qwen2.5 series,
designed for efficient
general-purpose language tasks.
It is offered in a Q4_K_M

 **qwen2.5:14b** 
qwen · ollama

Qwen 2.5 14B is a mid-size
model from Alibaba's Qwen 2.5
family, offering a strong balance
of reasoning and code
capabilities. It sits between the
7B and 32B variants, providing

 **qwen2.5:32b** 
qwen · ollama

Qwen2.5-32B is a
32-billion-parameter large
language model from the
Qwen2.5 series, built on the
transformer architecture
optimized for both text

 **qwen2.5:3b** 
qwen · ollama

Qwen2.5 3B is a compact,
instruction-tuned transformer
model from the Qwen2.5 series,
built to deliver high-quality text
generation while remaining
lightweight enough for on-device
reasoning (inference-optimized)

 **qwen2.5:72b** 
qwen · ollama

Qwen2.5-72B is a
72-billion-parameter
decoder-only transformer from
the Qwen2.5 series, released as
an open-source large language
model. It is optimized for both
code or resource-constrained

 **qwen2.5:7b** 
qwen · ollama

Qwen2.5-7B is a

 **qwen3-vl:32b** 
qwen · ollama

Qwen3-VL 32B is a

7-billion-parameter decoder-only
transformer from the Qwen2.5...

Tools LLM

Context Length: 32k



qwq:32b
qwen · ollama



qwq:32b is a
32-billion-parameter transform...

Thinking Tools LLM


Context Length: 32k

transformer-based multimodal
large language model from...

Vision Thinking Tools LLM

Context Length: 32k

📊 Most Popular Models (by token usage)





This project was developed with support from the National
Science Foundation (

NSF Award
#2427549

)